Modelling censored log-normal responses in generalized additive models

David L Miller (BioSS/UKCEH)

March 23, 2026

This document details the technical issues surrounding the implementation of censored log-normal responses in the R package `mgcv`. Necessary background information is provided along with derivations necessary for computation.

1 log-normal modelling

What does it mean to model log-normal data? What we are saying is that we have some variable, Y, which is $\log_b$-normally[1] distributed so $Y \sim \log_b N(\mu, \sigma^2)$ where μ is the logarithm of the mean and σ^2 is the base b logarithm of the variance of Y. This is equivalent to defining a variable $Z = b^Y$ where $Z \sim N(\mu_*, \sigma_*^2)$. The log-normal distribution is useful in cases where Y is defined over the interval $(0, \infty]$ (i.e., $Y > 0$).

Canonical references for the log-normal distribution include Aitchison and Brown (1957) and Crow and Shimizu (1988).

1.1 Arithmetic versus geometric means

We can estimate estimate the expectation of $\mathbb{E}(Y)$ by either an *arithmetic* or *geometric* mean (denoted $\mathbb{E}_A(Y)$ or $\mathbb{E}_G(Y)$, respectively). These are defined as:

$$\begin{aligned} \mathbb{E}_A(Y) &= b^{\mu + \frac{1}{2}\sigma^2} \quad \text{and} \\ \mathbb{E}_G(Y) &= b^{\mu}. \end{aligned} \tag{1}$$

Where the geometric "mean" corresponds to the median of the distribution. This difference is shown in Figure 1.

Similarly, we define the arithmetic and geometric standard deviations as follows:

$$\begin{aligned} \mathrm{SD}_A(Y) &= b^{\mu + \frac{1}{2}\sigma^2} \sqrt{b^{\sigma^2} - 1} \\ \mathrm{SD}_G(Y) &= b^{\sigma}. \end{aligned}$$

Following usual results, we have that multiplication of log-normal distributed variables results in a log-normally distributed product and therefore geometric means of log-normally distributed variables are log-normally distributed. We also have a multiplicative central limit theorem, where geometric means of (any) random variables are approximately log-normally distributed (under usual conditions).

1.2 log-normal response in regression-like models

We may want to model the mean via some linear model, for example when using a GLM or GAM[2]. We decompose the mean as $\mathbf{X}\boldsymbol{\beta}$, where $\mathbf{X}$ is an $n \times p$ design matrix and $\boldsymbol{\beta}$ is a p-vector of coefficients to be estimated (where n is the number of observations).

[1] Note that throughout I use $\log_a$ to indicate logarithm to the base a since these calculations work regardless of the base used. When the distribution is referred to informally (i.e., non-mathematically) I drop the base.

[2] Basic theory for log-normal linear models is covered in Section 3.2 of Crow and Shimizu (1988).

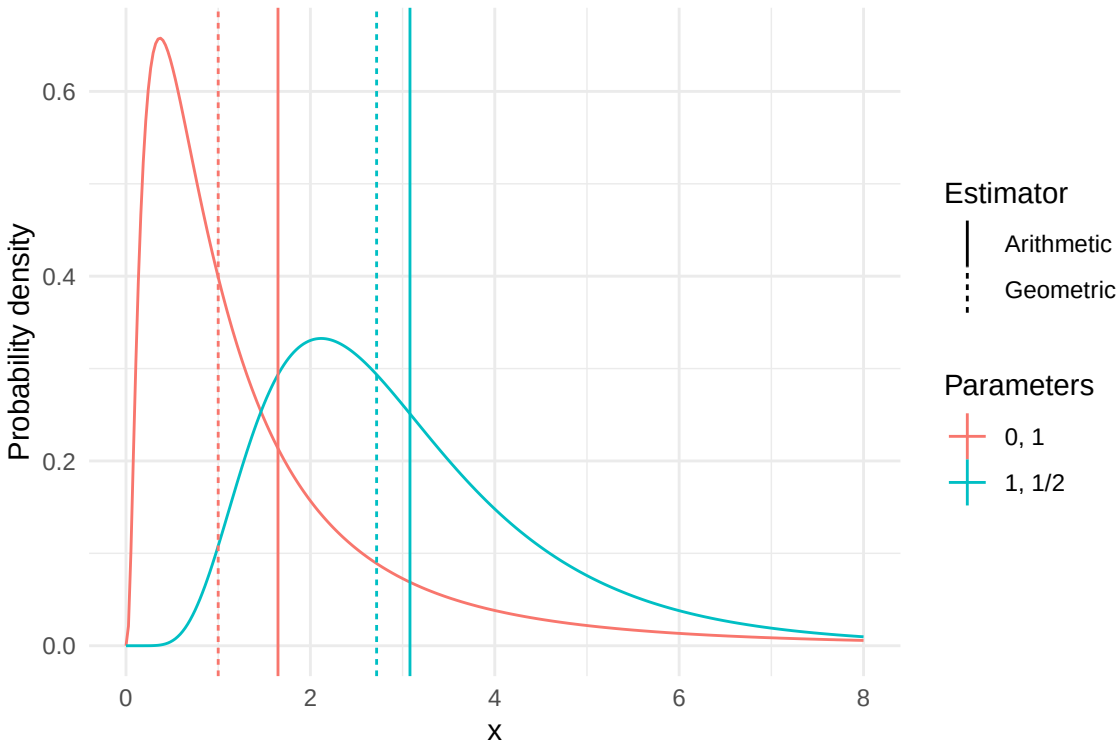

Figure 1: Demonstration of the differences in estimation of arithmetic vs geometric means for log-normal, overlaid on the distribution functions. Parameters are given on the $\log_e$ scale, as required by the R function `dlnorm`.

In practice, one can simply assume that the log-response is normally distributed and fit a GLM-a-like[3] with a normal distribution. This model is:

$$\mathbb{E}(Y) = b^{\mu} = b^{\mathbf{X}\boldsymbol{\beta}},$$

which is clearly modelling the geometric, rather than arithmetic mean (equivalently, we are modelling the median). This naturally extends to the case of random effects and therefore penalized smoothers (i.e., GLMMs, GAMs and GAMMs).

An example of such an analysis in R is as follows:

```
library(mgcv)

## Loading required package:  nlme
## This is mgcv 1.9-4.  For overview type '?mgcv'.

set.seed(2)

# generate some simple data
n <- 10000
x <- rnorm(n, sd=0.05)
beta0 <- 1
beta1 <- 1.2
mu <- beta0+beta1*x
s <- 0.2

# generate response
```

[3]GAM/GLMM/GAMM/etc.

```
z <- rlnorm(n, meanlog=mu, sdlog=s)

# fit model
aa <- gam(log(z)~x, method="REML")
summary(aa)
##
## Family: gaussian
## Link function: identity
##
## Formula:
## log(z) ~ x
##
## Parametric coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 0.999854   0.002024  493.92   <2e-16 ***
## x           1.255988   0.040499   31.01   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## R-sq.(adj) =  0.0877   Deviance explained = 8.78%
## -REML = -1778.1  Scale est. = 0.040974  n = 10000
# compare RSS
sum((z-exp(predict(aa)))^2)
## [1] 3274.107
sum((z-exp(predict(aa)+1/2*aa$sig2))^2)
## [1] 3242.197
```

2 Censored normal modelling

Censored normal models are one kind of tobit model[4] (Goldberger, 1964) where the response is censored in some way (in our case, that observations are below some detection threshold). Core to their formulation is the separation of the observation into censored and non-censored groups, where the censored likelihood contributions are evaluations of the cumulative distribution function (CDF) of the normal distribution.

Data takes the form of a triplet: $(y, \mathbf{x}, c)$ where y is the response, $\mathbf{x}$ is the corresponding covariate vector (or in the case of a GAM, vector of basis function evaluations), and c is an indicator which is 1 if y was censored at some point C (i.e., $y \leq C$) and 1 otherwise. Since we do not know the value of y when $y \leq C$, we need to use the cumulative density function to evaluate the likelihood up to C. For $y > C$ we can use the usual likelihood function, since we then know y.

For modelling censored normal variables, we conveniently know a lot about the CDF of the normal distribution; evaluation is therefore very simple and can use built-in functions in R.

2.1 Implementation in `mgcv` via `cnorm()`

From the `mgcv` help file `?cnorm`: "If y is the response with mean m and standard deviation $w^{-1/2}\exp(\theta)$, then $w^{1/2}(y-m)\exp(-\theta)$ follows an $N(0,1)$ distribution. That is $y \sim N(m, \exp(2\theta)/w)$. θ is a single scalar for all observations. Observations may be left, interval or right censored or uncensored."

[4] "Tobin's probit" tends to refer to any censored regression model.

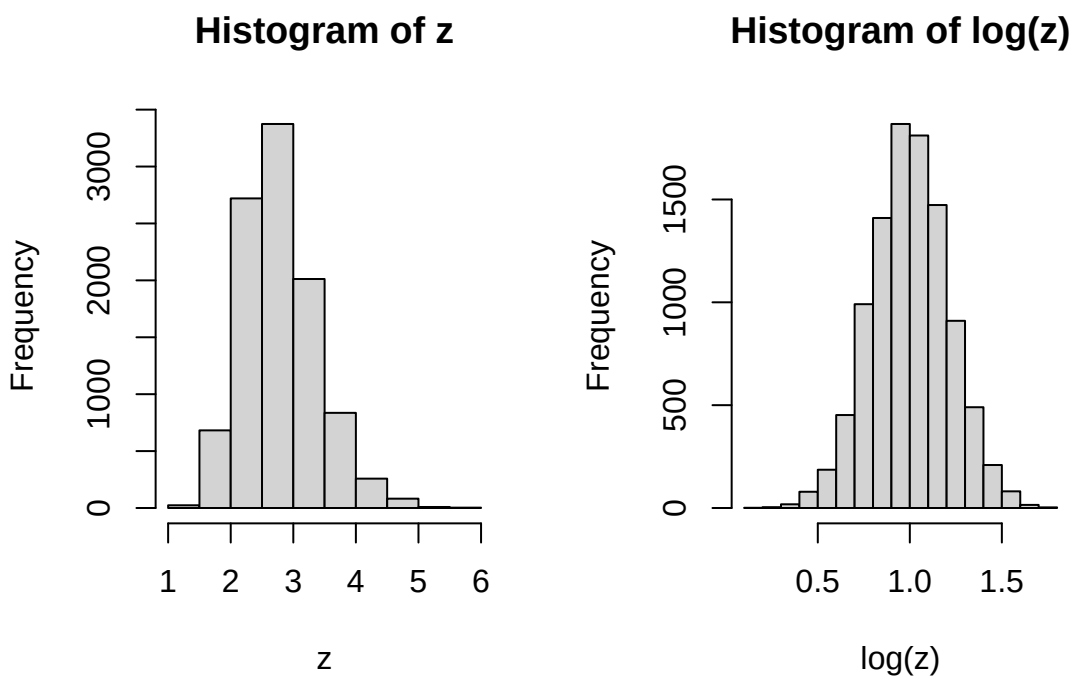

Figure 2: Histograms of the data generated in the linear model example. Left showing the data and right showing it log-transformed, looking more normal.

3 Combining approaches: censored log-normal distribution with arithmetic means for GAMs

We can combine these approaches using a geometric mean by simply combining the approach for modelling log-normal data described in Section (1.2) with the use of the censored normal distribution. In practice, that merely consists of using `mgcv::cnorm` with the log-normally-distributed data.

We can post hoc correct predictions to use the arithmetic mean using a multivariate version of the correction in (1). These corrections lead to the following expressions:

$$\begin{aligned}
\mathbb{E}\left[g\left(\hat{\boldsymbol{\beta}}\right)\right] &= g\left(\hat{\boldsymbol{\beta}}\right) + \frac{1}{2n}\text{trace}\left(D^2 g(\hat{\boldsymbol{\beta}})\Sigma\right) \quad \text{and} \\
\hat{\text{Var}}\left[g\left(\hat{\boldsymbol{\beta}}\right)\right] &= \frac{1}{n}\left[Dg(\boldsymbol{\beta})\right]^T \Sigma (Dg(\boldsymbol{\beta})) + \frac{1}{2n^2}\text{trace}\left(D^3 g(\boldsymbol{\beta}) D^2 g(\boldsymbol{\beta})\Sigma^2 - \left[D^2 g(\boldsymbol{\beta})\Sigma\right]^2\right)
\end{aligned} \tag{2}$$

where $g(\boldsymbol{\beta}) = b^{X\boldsymbol{\beta}}$ if b is the base we are working in and X is the design matrix for our model. D indicates the derivative operator with respect to $\boldsymbol{\beta}$, with powers of D giving that order of derivative. Σ is the posterior covariance matrix for $\boldsymbol{\beta}$ (see, e.g. Miller (2025) for further explanation of covariance matrices in the context of empirical Bayes modelling).

Derivations for these expressions are given in the Appendix.

4 Implementation of `mgcv` family: `clognorm()`

As the corrections in (2) can be applied post hoc, the easiest mechanism is via `mgcv`'s per-family prediction functions. Since version 1.8-0, it has been possible to implement a `$predict` function within a family (response distribution) function. When `mgcv::predict()` is used with the `type=''response''` flag set this function is called instead allowing

for per-family special prediction functions. It is this mechanism that we exploit to implement seamless predictions using base `mgcv` functions.

The `clognorm` family is therefore implemented in `mgcvUtils` as a copy of the `cnorm` library with the addition of a `$predict` method using (2). We allow an extra specification of the base to be used when calling the function during modelling (e.g., `base=10` for $\log_{10}$ or `base=''e''` for $\log_e$).

4.1 A quick example of usage

```
library(mgcv)
library(mgcvUtils)
library(ggplot2)

# make a 1D example with less noise
set.seed(42)

# simulated example from ?gam
dat <- gamSim(1, n=200, dist="normal", scale=1)
## Gu & Wahba 4 term additive model
e <- rnorm(nrow(dat), sd=2)
dat$y2 <- 10^(dat$f2/10 + e)

# censor values below some value
dat$y2 <- cbind(dat$y2, dat$y2)
# indicate which values are censored, see ?cnorm for rules
dat$y2[dat$y2[, 1] < 0.01, 2] <- -Inf
dat$y2[dat$y2[, 1] < 0.01, 1] <- 0.01

# response will be transformed by clognorm before fitting
b2 <- gam(y2~s(x2), data=dat, method="REML", family=clognorm(base=10))
summary(b2)
##
## Family: clog10norm(1.829)
## Link function: identity
##
## Formula:
## y2 ~ s(x2)
##
## Parametric coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.07974    0.13120   0.608    0.543
##
## Approximate significance of smooth terms:
##        edf Ref.df Chi.sq p-value
## s(x2) 4.28  5.278  11.12   0.058 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.0518   Deviance explained = 5.97%
## -REML = 382.63  Scale est. = 1         n = 200
```

We can compare the predictions from the two estimators in Figure (3).

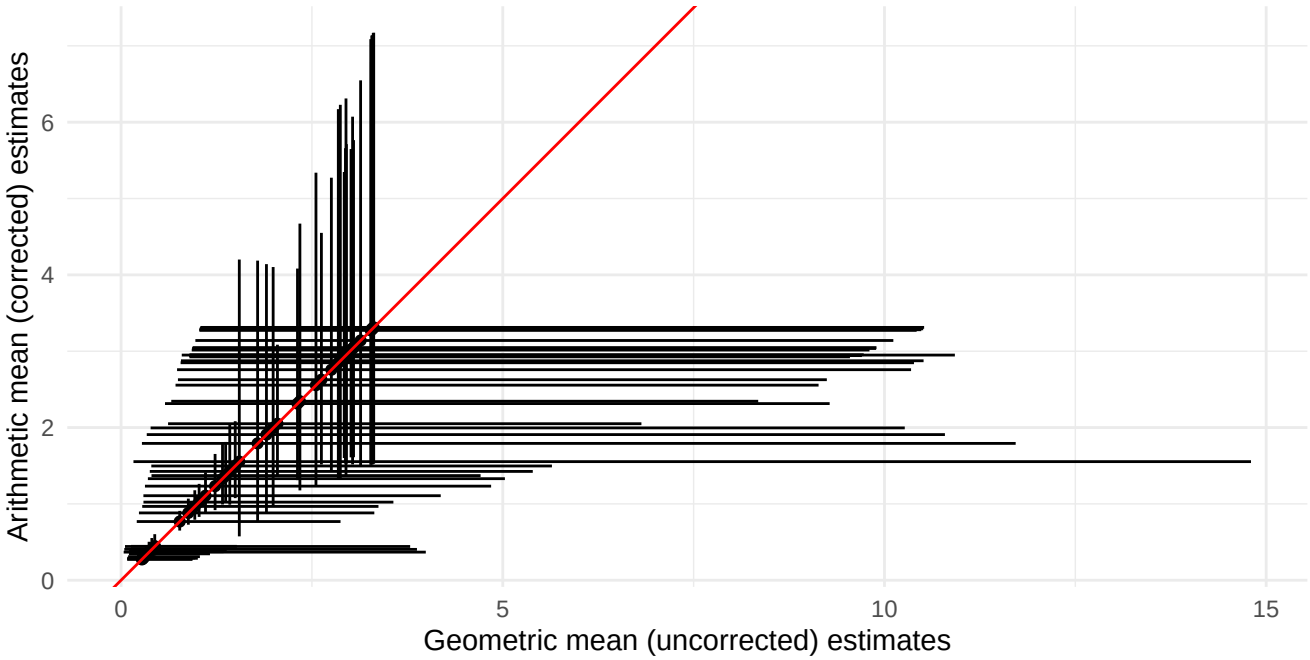

Figure 3: Comparison of the estimators.

Appendix

"Thorson's" epsilon correction

The core of the REML/ML methods implemented in `mgcv` is Laplace-approximate maximum likelihood (LAML). In the case where Laplace approximations are used, we can use the correction "proposed" by Thorson and Kristensen (2016) and elaborated on in Thorson and Kristensen (2024)[5]. Scare quotes are used here as the correction is based on Tierney et al. (1989), from their Theorem 2. They propose estimators for the expectation and variance of some function g of the model parameters ($\hat{\theta}$) when a model is estimated using a Laplace approximation.

We can generally write the expectation as:

$$\mathbb{E}\left[g(\theta)\right] = \frac{\int g(\theta)L(\theta)\pi(\theta)d\theta}{\int L(\theta)\pi(\theta)d\theta}. \tag{3}$$

As usual $L(\theta)$ is the likelihood and $\pi(\theta)$ is the prior on parameter θ. We can write (3) in the following form (equation 2.2 in Tierney et al. (1989)):

$$\mathbb{E}\left[g(\theta)\right] = \frac{\int b_N(\theta)\exp\left[-nh_N(\theta)\right]d\theta}{\int b_D(\theta)\exp\left[-nh_D(\theta)\right]d\theta}, \tag{4}$$

where N and D denote the numerator and denominator versions of the functions, respectively. We can see that our problem fits into the "standard form" where $b_N(\theta) = g(\theta)$, $b_D(\theta) = 1$ and $h_N(\theta) = h_D(\theta) = (1/n)\log_e\left[L(\theta)\pi(\theta)\right]$ (Tierney and Kadane, 1986), however Tierney et al. (1989) also introduce the "exponential form", which we will use below.

Analysis then proceeds by looking at $\mathbb{E}\left\{\exp[sg(\theta)]\right\}$ (the moment generating function). Plugging these values in and taking derivatives will yield Theorem 2 of Tierney et al. (1989) (see paper for fine details on this):

$$\begin{aligned}\mathbb{E}\left[g\left(\hat{\theta}\right)\right] &= g\left(\hat{\theta}\right) + \frac{d}{ds}\log_e\sigma_s|_{s=0} + \frac{d}{ds}\log_e b(\theta_s)|_{s=0} \quad \text{and} \\ \hat{\text{Var}}\left[g\left(\hat{\theta}\right)\right] &= \left[\sigma g'\left(\hat{\theta}\right)\right]^2/n + \frac{d^2}{ds^2}\log_e\sigma_s|_{s=0} + \frac{d^2}{ds^2}\log_e b(\theta_s)|_{s=0}.\end{aligned} \tag{5}$$

This is equivalent to expressions 2e and 2f in Thorson and Kristensen (2016) and equation 6.8 in Thorson and Kristensen (2024). Note that this is only for the case where we have parameter, θ, to estimate. In the case of Tierney et al. (1989), Thorson and Kristensen (2016) and Thorson and Kristensen (2024) the multivariate case is alluded to, but not fully expanded. The rest of the appendix derives this for the specific log-normal case.

Extending to the multivariate case

In the above, we are only looking at univariate θ, we need to extend to the multivariate case. Tierney et al. (1989) state (in the proof of Theorem 2):

"*In the multivariate statement of the theorem, σ_s would be replaced by $\det\left(\Sigma_s\right)^{1/2}$, where $\Sigma_s = \left[D^2h_s|_{\boldsymbol{\theta}_s}\right]$ and $\left[\sigma g'(\theta)\right]^2$ would be replaced by $(Dg)^T\Sigma(Dg)$, where Dg is the derivative of g [with respect to $\boldsymbol{\theta}$] evaluated at $\hat{\boldsymbol{\theta}}$ and $\Sigma = \left[D^2h|_{\hat{\boldsymbol{\theta}}}\right]$.*"

Note that here $h_s(\boldsymbol{\theta}) = nh(\boldsymbol{\theta}) - sg(\boldsymbol{\theta})$.

Doing the suggested substitution and letting $b = 1$, (5) becomes

$$\begin{aligned}\mathbb{E}\left[g\left(\hat{\boldsymbol{\theta}}\right)\right] &= g\left(\hat{\boldsymbol{\theta}}\right) + \frac{d}{ds}\log_e\det\left(\Sigma_s\right)^{1/2}|_{s=0} \quad \text{and} \\ \hat{\text{Var}}\left[g\left(\hat{\boldsymbol{\theta}}\right)\right] &= (Dg)^T\Sigma(Dg)/n + \frac{d^2}{ds^2}\log_e\det\left(\Sigma_s\right)^{1/2}\Big|_{s=0}.\end{aligned} \tag{6}$$

Note that the last term in each equation evaluates to zero since $b = 1$ (per above) and hence has no derivatives.

[5] Note that there are a couple of mathematical errors in both of these sources.

Finding $\frac{d\boldsymbol{\theta}_s}{ds}$

To find quantities such as $\frac{d}{ds}h(\boldsymbol{\theta}_s)$ we need to find $\frac{d\boldsymbol{\theta}_s}{ds}$ (since the resulting derivative is $h'(\boldsymbol{\theta}_s)\frac{d\boldsymbol{\theta}_s}{ds}$, by the chain rule). By analogy to equation 3.8 in Tierney et al. (1989), we find

$$\begin{aligned} 0 = \frac{\partial f}{\partial \boldsymbol{\theta}}|_{\boldsymbol{\theta}_s,s} &= \frac{\partial}{\partial \boldsymbol{\theta}}\left[nh(\boldsymbol{\theta}) - sg(\boldsymbol{\theta})\right]|_{\boldsymbol{\theta}_s,s} \\ &= n\frac{\partial}{\partial \boldsymbol{\theta}}h(\boldsymbol{\theta}) - s\frac{\partial}{\partial \boldsymbol{\theta}}g(\boldsymbol{\theta}). \end{aligned}$$

Then taking derivatives wrt s at $s = 0$, we have

$$\begin{aligned} 0 &= \frac{d}{ds}\left[n\frac{\partial}{\partial \boldsymbol{\theta}}h(\boldsymbol{\theta}) - s\frac{\partial}{\partial \boldsymbol{\theta}}g(\boldsymbol{\theta})\right]|_{s=0} \\ &= nD^2h(\hat{\boldsymbol{\theta}})\frac{d\boldsymbol{\theta}_s}{ds}|_{s=0} - \frac{\partial}{\partial \boldsymbol{\theta}}g(\boldsymbol{\theta}) \\ \frac{d\boldsymbol{\theta}_s}{ds}|_{s=0} &= \frac{1}{n}\frac{\partial}{\partial \boldsymbol{\theta}}g(\boldsymbol{\theta})\left[D^2h(\hat{\boldsymbol{\theta}})\right]^{-1} \\ &= \frac{1}{n}\frac{\partial}{\partial \boldsymbol{\theta}}g(\boldsymbol{\theta})\Sigma. \end{aligned}$$

Finding the derivatives of $\log_e \det(\Sigma_s)^{1/2}$

First derivatives are needed for the expectation:

$$
\begin{aligned}
\frac{d}{ds}\log_e\left[\det(\Sigma_s)^{1/2}\right]|_{s=0} &= \frac{\frac{d}{ds}\det(\Sigma_s)^{1/2}|_{s=0}}{\det(\Sigma)^{1/2}} \\
&= \frac{1/2\det(\Sigma)^{-1/2}\frac{d}{ds}\det\Sigma_s|_{s=0}}{\det(\Sigma)^{1/2}} \\
&= 1/2\det(\Sigma)^{-1}\frac{d}{ds}\det\Sigma_s|_{s=0} \\
&= 1/2\det(\Sigma)^{-1}\det\Sigma\mathrm{trace}\left(\Sigma^{-1}\frac{d}{ds}\Sigma_s|_{s=0}\right) \\
&= 1/2\mathrm{trace}\left(\Sigma^{-1}\frac{d}{ds}\Sigma_s|_{s=0}\right) \\
&= 1/2\mathrm{trace}\left(\Sigma^{-1}\frac{d}{ds}\left[yD^2h_s|_{\boldsymbol{\theta}_s}\right]^{-1}|_{s=0}\right) \\
&= -1/2\mathrm{trace}\left(\Sigma^{-1}\left[D^2h_s|_{\boldsymbol{\theta}_s}\right]^{-1}\frac{d}{ds}\left[D^2h(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}-\frac{s}{n}D^2g(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}\right]|_{s=0}\left[D^2h_s|_{\boldsymbol{\theta}_s}\right]^{-1}\right) \\
&= -1/2\mathrm{trace}\left(\Sigma^{-1}\Sigma\frac{d}{ds}\left[D^2h(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}-\frac{s}{n}D^2g(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}\right]|_{s=0}\Sigma\right) \\
&= -1/2\mathrm{trace}\left(\frac{d}{ds}\left[D^2h(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}-\frac{s}{n}D^2g(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}\right]|_{s=0}\Sigma\right) \\
&= -1/2\mathrm{trace}\left(\frac{d}{ds}\left[D^2h(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}-\frac{s}{n}D^2g(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}\right]|_{s=0}\Sigma\right) \\
&= -1/2\mathrm{trace}\left(\left[D^3h(\boldsymbol{\theta})\frac{d\boldsymbol{\theta}_s}{ds}-\frac{d}{ds}\left(\frac{s}{n}D^2g(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}\right)|_{s=0}\right]\Sigma\right)bg \\
&= -1/2\mathrm{trace}\left(\left[D^3h(\boldsymbol{\theta})\frac{1}{n}Dg(\boldsymbol{\theta})\Sigma-\frac{d}{ds}\left(\frac{s}{n}D^2g(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}\right)|_{s=0}\right]\Sigma\right) \\
&= -1/2\mathrm{trace}\left(\left[D^3h(\boldsymbol{\theta})\frac{1}{n}Dg(\boldsymbol{\theta})\Sigma-\left(\left(\frac{d}{ds}\frac{s}{n}|_{s=0}\right)D^2g(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}+\frac{s}{n}\left(\frac{d}{ds}D^2g(\boldsymbol{\theta})|_{s=0}\right)|_{\boldsymbol{\theta}_s}\right)\right]\Sigma\right) \\
&= -1/2\mathrm{trace}\left(\left[D^3h(\boldsymbol{\theta})\frac{1}{n}Dg(\boldsymbol{\theta})\Sigma-\left(\left(\frac{d}{ds}\frac{s}{n}|_{s=0}\right)D^2g(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}\right)\right]\Sigma\right) \\
&= -1/2\mathrm{trace}\left(\left[D^3h(\boldsymbol{\theta})\frac{1}{n}Dg(\boldsymbol{\theta})\Sigma-\left(\left(\frac{d}{ds}\frac{s}{n}|_{s=0}\right)D^2g(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}\right)\right]\Sigma\right) \\
&= -1/2\mathrm{trace}\left(\left[\frac{1}{n}D^3h(\boldsymbol{\theta})Dg(\boldsymbol{\theta})\Sigma-\frac{1}{n}D^2g(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}\right]\Sigma\right) \\
&= \frac{1}{2n}\mathrm{trace}\left(D^2g(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}\Sigma-D^3h(\boldsymbol{\theta})Dg(\boldsymbol{\theta})\Sigma^2\right)
\end{aligned}
$$

This is analogous to the expression given in equation 6.8 in Thorson and Kristensen (2024).

For the variance expression, we need to also evaluate $\frac{d^2}{ds^2}\log_e\det(\Sigma_s)^{1/2}$. To save on space (and further mathematical headaches) we assume in the below that we have a normal likelihood, such that the third derivatives of h do not exist. This allows us to drop terms we don't use later. This means that the following is not completely

general.

$$
\begin{aligned}
\frac{d^2}{ds^2}\log_e \det\left(\Sigma_s\right)^{1/2}|_{s=0} &= \frac{d}{ds}\left.\frac{\frac{d}{ds}\det\left(\Sigma_s\right)^{1/2}}{\det\left(\Sigma\right)^{1/2}}\right|_{s=0} \\
&= \frac{d}{ds}\det\left(\Sigma_s\right)^{-1/2}\frac{d}{ds}\det\left(\Sigma_s\right)^{1/2}|_{s=0} \\
&= \frac{d}{ds}\left[\frac{1}{2}\det\left(\Sigma_s\right)^{-1/2}\det\left(\Sigma_s\right)^{-1/2}\frac{d}{ds}\det\left(\Sigma_s\right)\right]|_{s=0} \\
&= \frac{1}{2}\frac{d}{ds}\left[\det\left(\Sigma_s\right)^{-1}\frac{d}{ds}\det\left(\Sigma_s\right)\right]|_{s=0} \\
&= \frac{1}{2}\frac{d}{ds}\left[\det\left(\Sigma_s\right)^{-1}\det\left(\Sigma_s\right)\operatorname{trace}\left(\Sigma_s^{-1}\frac{d}{ds}\Sigma_s\right)\right]|_{s=0} \\
&= \frac{1}{2}\frac{d}{ds}\left[\operatorname{trace}\left(\Sigma_s^{-1}\frac{d}{ds}\Sigma_s\right)\right]|_{s=0} \\
&= \frac{1}{2}\frac{d}{ds}\left[\operatorname{trace}\left(\Sigma_s^{-1}(-1)\left[D^2h_s|_{\boldsymbol{\theta}_s}\right]^{-1}\frac{d}{ds}\left[D^2h(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}-\frac{s}{n}D^2g(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}\right]\left[D^2h_s|_{\boldsymbol{\theta}_s}\right]^{-1}\right)\right]|_{s=0} \\
&= -\frac{1}{2}\frac{d}{ds}\left[\operatorname{trace}\left(\Sigma_s^{-1}\Sigma\frac{d}{ds}\left[D^2h(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}-\frac{s}{n}D^2g(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}\right]\Sigma\right)\right]|_{s=0} \\
&= -\frac{1}{2}\frac{d}{ds}\left[\operatorname{trace}\left(\frac{d}{ds}\left[D^2h(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}-\frac{s}{n}D^2g(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}\right]\Sigma\right)\right]|_{s=0} \\
&= -\frac{1}{2}\frac{d}{ds}\left[\operatorname{trace}\left(\left[\frac{d}{ds}D^2h(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}-\frac{d}{ds}\frac{s}{n}D^2g(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}\right]\Sigma\right)\right]|_{s=0} \\
&= -\frac{1}{2}\frac{d}{ds}\left[\operatorname{trace}\left(\left[-\frac{d}{ds}\frac{s}{n}D^2g(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}\right]\Sigma\right)\right]|_{s=0} \\
&= \frac{1}{2}\frac{d}{ds}\left[\operatorname{trace}\left(\frac{1}{n}D^2g(\boldsymbol{\theta})\Sigma\right)\right]|_{s=0} \\
&= \frac{1}{2n}\operatorname{trace}\left(\frac{d}{ds}D^2g(\boldsymbol{\theta})\Sigma_s|_{s=0}\right) \\
&= \frac{1}{2n}\operatorname{trace}\left(\left[\frac{d}{ds}D^2g(\boldsymbol{\theta})|_{s=0}\right]\Sigma+D^2g(\boldsymbol{\theta})\frac{d}{ds}\Sigma_s|_{s=0}\right) \\
&= \frac{1}{2n}\operatorname{trace}\left(\left[D^3g(\boldsymbol{\theta})\frac{1}{n}D^2g(\boldsymbol{\theta})\Sigma\right]\Sigma+D^2g(\boldsymbol{\theta})\Sigma\frac{d}{ds}\left[D^2h(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}-\frac{s}{n}D^2g(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}\right]|_{s=0}\Sigma\right) \\
&= \frac{1}{2n}\operatorname{trace}\left(\left[\frac{1}{n}D^3g(\boldsymbol{\theta})D^2g(\boldsymbol{\theta})\Sigma\right]\Sigma+D^2g(\boldsymbol{\theta})\Sigma\left[\frac{d}{ds}D^2h(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}-\frac{d}{ds}\frac{s}{n}D^2g(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}\right]|_{s=0}\Sigma\right) \\
&= \frac{1}{2n}\operatorname{trace}\left(\left[\frac{1}{n}D^3g(\boldsymbol{\theta})D^2g(\boldsymbol{\theta})\Sigma\right]\Sigma+D^2g(\boldsymbol{\theta})\Sigma\left[D^3h(\boldsymbol{\theta})\frac{d\boldsymbol{\theta}}{ds}|_{s=0}-\frac{1}{n}D^2g(\boldsymbol{\theta})\right]\Sigma\right) \\
&= \frac{1}{2n}\operatorname{trace}\left(\left[\frac{1}{n}D^3g(\boldsymbol{\theta})D^2g(\boldsymbol{\theta})\Sigma\right]\Sigma+D^2g(\boldsymbol{\theta})\Sigma\left[-\frac{1}{n}D^2g(\boldsymbol{\theta})\right]\Sigma\right) \\
&= \frac{1}{2n^2}\operatorname{trace}\left(D^3g(\boldsymbol{\theta})D^2g(\boldsymbol{\theta})\Sigma\Sigma-D^2g(\boldsymbol{\theta})\Sigma D^2g(\boldsymbol{\theta})\Sigma\right) \\
&= \frac{1}{2n^2}\operatorname{trace}\left(D^3g(\boldsymbol{\theta})D^2g(\boldsymbol{\theta})\Sigma^2-\left[D^2g(\boldsymbol{\theta})\Sigma\right]^2\right).
\end{aligned}
$$

Putting it together

Substituting the above into (6) we then obtain:

$$
\begin{aligned}
\mathbb{E}\left[g\left(\hat{\boldsymbol{\theta}}\right)\right] &= g\left(\hat{\boldsymbol{\theta}}\right) + \frac{1}{2n}\text{trace}\left(D^2 g(\boldsymbol{\theta})|_{\boldsymbol{\theta}_s}\Sigma - D^3 h(\boldsymbol{\theta}) Dg(\boldsymbol{\theta})\Sigma^2\right) \quad \text{and} \\
\hat{\text{Var}}\left[g\left(\hat{\boldsymbol{\theta}}\right)\right] &= (Dg)^T\Sigma(Dg)/n + \left.\frac{d^2}{ds^2}\log_e \det\left(\Sigma_s\right)^{1/2}\right|_{s=0}.
\end{aligned} \tag{7}
$$

Implications for log-normal modelling

When we model log-normal data, we are letting $h(\theta)$ be a normal distribution pdf. In that case, we know that the third derivative of the normal distribution with respect to θ (the parameters of the linear predictor) is 0, so the second term inside the trace() in the expectation in (7) will cancel. We therefore obtain the following corrections:

$$
\begin{aligned}
\mathbb{E}\left[g\left(\hat{\boldsymbol{\theta}}\right)\right] &= g\left(\hat{\boldsymbol{\theta}}\right) + \frac{1}{2n}\text{trace}\left(D^2 g(\hat{\boldsymbol{\theta}})\Sigma\right) \quad \text{and} \\
\hat{\text{Var}}\left[g\left(\hat{\boldsymbol{\theta}}\right)\right] &= (Dg)^T\Sigma(Dg)/n + \frac{1}{2n^2}\text{trace}\left(D^3 g(\boldsymbol{\theta}) D^2 g(\boldsymbol{\theta})\Sigma^2 - \left[D^2 g(\boldsymbol{\theta})\Sigma\right]^2\right).
\end{aligned} \tag{8}
$$

In practice, we can obtain Σ as the posterior covariance matrix (via `mgcv::vcov`, see Wood et al. (2016) and Miller (2025) for more information on the meaning of these covariance matrices and how to include uncertainty from smoothing parameters).

We can obtain $D^2 g(\hat{\boldsymbol{\theta}}) = \left.\frac{\partial^2 g(\boldsymbol{\theta})}{\partial \boldsymbol{\theta}}\right|_{\boldsymbol{\theta}=\hat{\boldsymbol{\theta}}}$ (i.e., the hessian of $g(\boldsymbol{\theta})$ at the REML-optimal parameters) by numerical derivatives (e.g., `numDeriv::Hessian`; Gilbert and Varadhan (2006)) or by analytical expression. Since the analytical expression is faster to compute (since we know that $g(\boldsymbol{\beta}) = b^{X\boldsymbol{\beta}}$), we calculate the j, k^{th} entry in the hessian is given by:

$$
\begin{aligned}
\frac{\partial^2 b^{X\boldsymbol{\beta}}}{\partial\beta_j\partial\beta_k} &= \frac{\partial}{\partial\beta_j}\frac{\partial}{\partial\beta_k} b^{X\boldsymbol{\beta}} \\
&= \frac{\partial}{\partial\beta_j}\left[\left(\frac{\partial}{\partial\beta_k} X\boldsymbol{\beta}\right) b^{X\boldsymbol{\beta}}\log_e b\right] \\
&= \frac{\partial}{\partial\beta_j}\left(\boldsymbol{x}_k^T b^{X\boldsymbol{\beta}}\log_e b\right) \\
&= \boldsymbol{x}_k^T\left(\log_e b\right)\frac{\partial}{\partial\beta_j} b^{X\boldsymbol{\beta}} \\
&= \boldsymbol{x}_k^T\left(\log_e b\right)\boldsymbol{x}_j b^{X\boldsymbol{\beta}}\left(\log_e b\right) \\
&= \left(\log_e b\right)^2 \boldsymbol{x}_k^T\boldsymbol{x}_j b^{X\boldsymbol{\beta}}
\end{aligned}
$$

where b is the base of the response (i.e., we are modelling on the $\log_b$ scale) and $\boldsymbol{x}_j$ and $\boldsymbol{x}_k$ indicate a rows/columns of the design matrix X. By extension (and because we need it for the variance term):

$$
\frac{\partial^3 b^{X\boldsymbol{\beta}}}{\partial\beta_j\partial\beta_k\partial\beta_l} = \left(\log_e b\right)^3\left(\boldsymbol{x}_l \times \boldsymbol{x}_k^T\boldsymbol{x}_j\right) b^{X\boldsymbol{\beta}},
$$

where $\left(\boldsymbol{x}_l \times \boldsymbol{x}_k^T\boldsymbol{x}_j\right)$ is an order 3 tensor (i.e., a 3-dimensional array).

Dealing with order 3 tensors

Tensor algebra is tricky and can cause a headache, especially when performing operations on a computer. To make things easier (and so we only need use matrix algebra) we can "unpack" our 3-tensor to a matrix and then using Kronecker products. We first need to remember that the tensor-matrix product will be necessarily dimension-reducing, by analogy to matrix-vector multiplication giving a vector result. Hence tensor-matrix product gives a matrix result.

If we have a rows-by-columns-by-slices 3-tensor, we can "unpack" that into a rows-by-(columns-times-slices) matrix instead and then any multiplication that was to be on that tensor can be made on the unpacked version, if we can create an appropriate matrix to multiply it by. The "appropriate" operation is the Kronecker product, which takes two matrices and multiplies the whole of one matrix by each element of the other, expanding it in both rows and columns:

$$A \otimes B = \begin{pmatrix} A_{11}B & A_{12}B & \dots & A_{1n}B \\ A_{21}B & A_{22}B & & A_{2n}B \\ \vdots & & \ddots & \vdots \\ A_{n1}B & A_{n2}B & \dots & A_{nn}B \end{pmatrix},$$

where $\otimes$ denotes the Kronecker product and A_{ij} indicates the i, j^{th} entry in the matrix A.

You want to multiply a 3-tensor T by some matrix M. Take the slices of T, and lay them out to make a wide matrix T^* so we have $T^* = (T_{..1}, T_{..2}, \dots)$ if $T_{..i}$ indicates the ith slice of T. Then form the Kronecker product of M with the identity matrix that has the same size as the number of slices, $M \otimes I$. The crossproduct of T^* and the Kronecker product, $T^*(M \otimes I)$ is then equivalent to the product TM: $T^*(M \otimes I) = TM$. This can be verified by writing out what the elements of each side of this equality are.

See Kolda and Bader (2009) for more information on tensors and multilinear algebra.

References

J. J. Aitchison and J. Brown. *The Lognormal Distribution, with Special Reference to Its Uses in Economics*. Number 5 in University of Cambridge Department of Economics Monograph. Cambridge, University Press, 1957. ISBN 978-0-521-04011-2.

E. L. Crow and K. Shimizu, editors. *Lognormal Distributions: Theory and Applications*. Number vol. 88 in Statistics, Textbooks and Monographs. M. Dekker, New York, 1988. ISBN 978-0-8247-7803-3.

P. Gilbert and R. Varadhan. numDeriv: Accurate Numerical Derivatives, Apr. 2006.

A. S. Goldberger. *Econometric Theory*. J. Wiley, 1964. ISBN 978-0-471-31101-0.

T. G. Kolda and B. W. Bader. Tensor Decompositions and Applications. *SIAM Review*, 51(3):455–500, Aug. 2009. ISSN 0036-1445, 1095-7200. doi: 10.1137/07070111X.

D. L. Miller. Bayesian views of generalized additive modelling. *Methods in Ecology and Evolution*, 16(3):446–455, 2025. ISSN 2041-210X. doi: 10.1111/2041-210X.14498.

J. Thorson and K. Kristensen. *Spatio-Temporal Models for Ecologists*. Chapman and Hall/CRC, New York, Feb. 2024. ISBN 978-1-003-41029-4. doi: 10.1201/9781003410294.

J. T. Thorson and K. Kristensen. Implementing a generic method for bias correction in statistical models using random effects, with spatial and population dynamics examples. *Fisheries Research*, 175:66–74, Mar. 2016. ISSN 01657836. doi: 10.1016/j.fishres.2015.11.016.

L. Tierney and J. B. Kadane. Accurate Approximations for Posterior Moments and Marginal Densities. *Journal of the American Statistical Association*, 81(393):82–86, Mar. 1986. ISSN 0162-1459. doi: 10.1080/01621459.1986.10478240.

L. Tierney, R. E. Kass, and J. B. Kadane. Fully Exponential Laplace Approximations to Expectations and Variances of Nonpositive Functions. *Journal of the American Statistical Association*, 84(407):710–716, 1989.

S. N. Wood, N. Pya, and B. Säfken. Smoothing Parameter and Model Selection for General Smooth Models. *Journal of the American Statistical Association*, 111(516):1548–1563, Oct. 2016. ISSN 0162-1459, 1537-274X. doi: 10.1080/01621459.2016.1180986.